

## ▼ 중고차 가격 예측

데이터 출처 : <https://www.kaggle.com/datasets/adityadesai13/used-car-dataset-ford-and-mercedes?select=vw.csv>

수정 데이터 출처 : [https://www.datamanim.com/dataset/03\\_dataq/typetwo.html#id16](https://www.datamanim.com/dataset/03_dataq/typetwo.html#id16)

x\_train: [https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/X\\_train.csv](https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/X_train.csv)

y\_train: [https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/y\\_train.csv](https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/y_train.csv)

x\_test: [https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/X\\_test.csv](https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/X_test.csv)

x\_label(평가용) :

[https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/y\\_test.csv](https://raw.githubusercontent.com/Datamanim/datarepo/main/carsprice/y_test.csv)

project file(.ipynb) : <https://github.com/fa-ina-tic/report/blob/main/UsedCar.ipynb>

```
import pandas as pd
#데이터 로드
x_train = pd.read_csv("https://raw.githubusercontent.com/Datamanim/datarepo/main/ca
y_train = pd.read_csv("https://raw.githubusercontent.com/Datamanim/datarepo/main/ca
x_test= pd.read_csv("https://raw.githubusercontent.com/Datamanim/datarepo/main/cars

#import data 확인
display(x_train.head())
display(y_train.head())
```

|   | carID | brand    | model    | year | transmission | mileage | fuelType | tax   | mpg  | engi |
|---|-------|----------|----------|------|--------------|---------|----------|-------|------|------|
| 0 | 13207 | hyundi   | Santa Fe | 2019 | Semi-Auto    | 4223    | Diesel   | 145.0 | 39.8 |      |
| 1 | 17314 | vauxhall | GTC      | 2015 | Manual       | 47870   | Diesel   | 125.0 | 60.1 |      |
| 2 | 12342 | audi     | RS4      | 2019 | Automatic    | 5151    | Petrol   | 145.0 | 29.1 |      |
| 3 | 13426 | vw       | Scirocco | 2016 | Automatic    | 20423   | Diesel   | 30.0  | 57.6 |      |
| 4 | 16004 | skoda    | Scala    | 2020 | Semi-Auto    | 3569    | Petrol   | 145.0 | 47.1 |      |

|   | carID | price |
|---|-------|-------|
| 0 | 13207 | 31995 |
| 1 | 17314 | 7700  |
| 2 | 12342 | 58990 |
| 3 | 13426 | 12999 |
| 4 | 16004 | 16990 |

## ▼ EDA

### 1. 결측값

## 2. 이상값

## 3. 데이터 타입

```
# data 기본 정보 확인
# X : 9 Columns 4960 Rows Null = None
# Y : 2 Columns 4960 Rows Null = None
print(x_train.info())
print(y_train.info())
print(x_test.info())

print(x_train.describe())
print(y_train.describe())
print(x_test.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4960 entries, 0 to 4959
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   carID            4960 non-null   int64
1   brand            4960 non-null   object
2   model            4960 non-null   object
3   year             4960 non-null   int64
4   transmission     4960 non-null   object
5   mileage          4960 non-null   int64
6   fuelType         4960 non-null   object
7   tax              4960 non-null   float64
8   mpg              4960 non-null   float64
9   engineSize       4960 non-null   float64
dtypes: float64(3), int64(3), object(4)
memory usage: 387.6+ KB
None
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4960 entries, 0 to 4959
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   carID   4960 non-null   int64
1   price   4960 non-null   int64
dtypes: int64(2)
memory usage: 77.6 KB
None
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2672 entries, 0 to 2671
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   carID            2672 non-null   int64
1   brand            2672 non-null   object
2   model            2672 non-null   object
3   year             2672 non-null   int64
4   transmission     2672 non-null   object
5   mileage          2672 non-null   int64
6   fuelType         2672 non-null   object
7   tax              2672 non-null   float64
8   mpg              2672 non-null   float64
```

```

9    engineSize    2672 non-null    float64
dtypes: float64(3), int64(3), object(4)
memory usage: 208.9+ KB
None

```

|       | carID        | year        | mileage       | tax         | mpg \       |
|-------|--------------|-------------|---------------|-------------|-------------|
| count | 4960.000000  | 4960.000000 | 4960.000000   | 4960.000000 | 4960.000000 |
| mean  | 15832.446169 | 2016.737903 | 24956.286895  | 152.332661  | 50.370766   |
| std   | 2206.717006  | 2.884035    | 24443.333662  | 82.403844   | 35.746505   |
| min   | 12002.000000 | 1997.000000 | 1.000000      | 0.000000    | 2.800000    |
| 25%   | 13929.250000 | 2016.000000 | 5641.250000   | 145.000000  | 38.700000   |
| 50%   | 15840.000000 | 2017.000000 | 19000.000000  | 145.000000  | 47.100000   |
| 75%   | 17765.750000 | 2019.000000 | 36702.000000  | 150.000000  | 54.300000   |
| max   | 19629.000000 | 2020.000000 | 259000.000000 | 580.000000  | 470.800000  |

```

engineSize
count    4960.000000

```

# 결측값

```

print(x_train.isnull().sum())
print(x_train.isnull().sum())
print(x_train.isnull().sum())

```

```

carID      0
brand      0
model      0
year       0
transmission 0
mileage    0
fuelType   0
tax        0
mpg        0
engineSize 0
dtype: int64
carID      0
brand      0
model      0
year       0
transmission 0
mileage    0
fuelType   0
tax        0
mpg        0
engineSize 0
dtype: int64
carID      0
brand      0
model      0
year       0
transmission 0
mileage    0
fuelType   0
tax        0
mpg        0
engineSize 0
dtype: int64

```

# 이상값

```
# 이상값 제거 함수 정의
def get_outliers(df=None, column=None, weight=1.5):
    per_75 = np.percentile(df[column].values, 75)
    per_25 = np.percentile(df[column].values, 25)


    IQR = (per_75 - per_25) * weight
    high = per_75 + IQR
    low = per_25 - IQR

    outlier_idx = df[(df[column]>high)|(df[column]<low)].index
    return outlier_idx

import matplotlib.pyplot as plt
import seaborn as sns





# 시각화로 탐색
# 1. 이상치 개수 파악
boxplot = sns.boxplot(data=x_train[['mileage', 'tax', 'mpg', 'engineSize']])
plt.show()

boxplot = sns.boxplot(data=x_train[['tax', 'mpg', 'engineSize']])
plt.show()
```



```
# 2. 값들의 분포 확인
for i in ['mileage', 'tax', 'mpg', 'engineSize']:
    plt.hist(x = x_train[[i]])
```

```
plt.show()
```



```
# mileage, mpg column에 로그를 취하여 머신 러닝의 성능을 높인다.
```

```
import numpy as np
```

```
log_features = ['mileage', 'mpg']
```

```
for i in log_features:
```



```
    x_train[i] = x_train[i].apply(lambda x: np.log1p(x))
```

```
    x_test[i] = x_test[i].apply(lambda x: np.log1p(x))
```

```
for i in log_features:
```

```
    plt.hist(x = x_train[[i]])
```

```
    plt.show()
```



```
for i in ['mileage', 'tax', 'mpg', 'engineSize']:
```

```
    outliers = get_outliers(x_train, i, 2)
```

```
    x_train.drop(outliers, axis=0, inplace=True)
```

```
    y_train.drop(outliers, axis=0, inplace=True)
```

```
# data type
```

```
# 명목변수 분리
```

```
cate_col = ['brand', 'model', 'year', 'transmission', 'fuelType']
```

```
# string / float&int
string_col = ['brand', 'model', 'transmission', 'fuelType']
num_col = list(x_train.columns.drop(string_col))
```

```
# 명목 변수 data type category로 변경
for i in cate_col:
    x_train[i] = x_train[i].astype('category')
    x_test[i] = x_test[i].astype('category')
print(x_train.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3070 entries, 0 to 4959
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   carID                 3070 non-null   int64
1   brand                 3070 non-null   category
2   model                 3070 non-null   category
3   year                  3070 non-null   category
4   transmission          3070 non-null   category
5   mileage               3070 non-null   float64
6   fuelType              3070 non-null   category
7   tax                   3070 non-null   float64
8   mpg                   3070 non-null   float64
9   engineSize            3070 non-null   float64
dtypes: category(5), float64(4), int64(1)
memory usage: 163.0 KB
None
```

```
# train과 test에 명목 변수가 다 같이 있는지 확인
train_dum = pd.get_dummies(x_train[cate_col])
test_dum = pd.get_dummies(x_test[cate_col])
```

```
print(set(train_dum.columns) - set(test_dum.columns))
print(set(test_dum.columns) - set(train_dum.columns))
```

```
{'model_ M6'}
{'model_ IQ', 'model_ Amarok', 'model_ CLK', 'model_ Hilux', 'model_ Californi
```

```
# 없는 명목변수를 포함하는 train, test 테이블 생성
X_train_dum = pd.get_dummies(x_train)
X_test_dum = pd.get_dummies(x_test)
```

```
train_missing = set(test_dum.columns) - set(train_dum.columns)
test_missing = set(train_dum.columns) - set(test_dum.columns)
```

```
for c in train_missing:
    X_train_dum[c] = 0
for c in test_missing:
    X_test_dum[c] = 0
```

## ▼ Model Selection

```
# package import
from sklearn.model_selection import KFold, train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from xgboost import XGBRegressor

X = X_train_dum.iloc[:, 1:]
y = y_train.iloc[:, 1:]

kfold = KFold(n_splits=5, shuffle=True)
pipe = Pipeline(steps=[('model', XGBRegressor(eval_metric='mlogloss', use_label_enc
grid_params = {'model':[XGBRegressor(eval_metric='mlogloss', use_label_encoder=False,
    'model__max_depth':[3, 5, 7],
    'model__learning_rate':[0.05, 0.1, 0.2, 0.3],
    'model__n_estimators':[50, 100, 200, 300],
    'model__gamma':[0, 0.1, 0.2]}

grid = GridSearchCV(pipe, grid_params, cv=kfold, verbose=0)
grid.fit(X, y)
print(grid.best_params_)
print(grid.best_score_)

{'model': XGBRegressor(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric='mlogloss', gamma=0,
    gpu_id=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=0.2, max_bin=None,
    max_cat_to_onehot=None, max_delta_step=None, max_depth=5,
    max_leaves=None, min_child_weight=None, missing=nan,
    monotone_constraints=None, n_estimators=300, n_jobs=None,
    num_parallel_tree=None, predictor=None, random_state=None,
    reg_alpha=None, reg_lambda=None, ...), 'model__gamma': 0, 'model_
0.9612029380888971
```

## ▼ Modeling

```
model = XGBRegressor(eval_metric='mlogloss', use_label_encoder=False,
    learning_rate=0.2, max_depth=5, n_estimators=200)
model.fit(X, y)
pred = model.predict(X_test_dum.iloc[:, 1:])
pred = pd.Series(pred)

answer = pd.concat([X_test_dum.carID, pred], axis=1)

answer
```



|             | <b>carID</b> | <b>0</b>     |
|-------------|--------------|--------------|
| <b>0</b>    | 12000        | 38508.203125 |
| <b>1</b>    | 12001        | 33022.828125 |
| <b>2</b>    | 12004        | 50232.843750 |
| <b>3</b>    | 12013        | 17864.957031 |
| <b>4</b>    | 12017        | 80702.773438 |
| ...         | ...          | ...          |
| <b>2667</b> | 19618        | 72973.968750 |
| <b>2668</b> | 19620        | 16821.384766 |
| <b>2669</b> | 19626        | 16913.240234 |
| <b>2670</b> | 19630        | 25904.589844 |

```
x_label = pd.read_csv("https://raw.githubusercontent.com/Datamanim/datarepo/main/ca")
display(x_label.head())
```

|          | <b>carID</b> | <b>price</b> |
|----------|--------------|--------------|
| <b>0</b> | 12000        | 38000        |
| <b>1</b> | 12001        | 23495        |
| <b>2</b> | 12004        | 59999        |
| <b>3</b> | 12013        | 16713        |
| <b>4</b> | 12017        | 46000        |

```
from sklearn.metrics import r2_score

score = r2_score(x_label.price, answer.iloc[:, 1])
print(score)

0.7104951468890918
```

Colab 유료 제품 - [여기에서 계약 취소](#)

